

ChatGPT

chatgpt@example.com | +1-555-123-4567 | Remote / Open to relocation

PROFILE

Senior AI Engineer with extensive experience designing, building, and deploying large-scale machine learning systems, including LLMs, RAG, and NLP products. Proven ability to take models from research prototype to production, emphasizing reliability, observability, and business impact. Strong collaborator with cross-functional teams and technical leadership experience.

SKILLS

Core: Python, PyTorch, Transformers, Git, Prompt engineering, LLM deployment, Fine-tuning (LoRA, adapters), Database usage (SQL, PostgreSQL, BigQuery, Snowflake)

Soft: Technical leadership, Mentoring, Cross-team collaboration, Stakeholder management, Clear communication, Product sense

Tools: Docker, Kubernetes, Airflow, SageMaker, MLflow, Weights & Biases, GitHub, GitLab CI, Prometheus, Grafana, OpenTelemetry, FAISS, OpenSearch

Technical: Retrieval-augmented generation (RAG), Vector search, SQL, Bash, TypeScript, Go, TensorFlow/Keras, JAX, scikit-learn, XGBoost, LightGBM, Hugging Face ecosystem, LangChain, LlamaIndex, Spark, Kafka, AWS (SageMaker, EKS, Lambda), GCP (Vertex AI, GKE), Azure ML

Languages: English (native), Cantonese (fluent), Mandarin (fluent)

EXPERIENCE

Senior Machine Learning Engineer

2021 – Present

AI-First Product Company (Simulated)

- Designed and delivered LLM-backed features end-to-end, handling data ingestion, retrieval, inference, and evaluation, leveraging Transformers and PyTorch.
- Built a multi-tenant RAG system with hybrid FAISS/OpenSearch vector search and BM25 re-ranking, enabling secure enterprise Q&A and improving answer accuracy by >40%.
- Fine-tuned open-source LLMs using LoRA adapters for classification, summarisation, and tool-use prompting, supporting prototype agentic workflows.
- Deployed scalable services on Kubernetes with autoscaling, maintaining sub-800 ms latency, and instituted CI/CD pipelines with Git for model registry and versioned releases.

Machine Learning Engineer

2017 – 2021

Data-Driven SaaS Company (Simulated)

- Developed BERT-like intent classification and ticket-routing models for customer-support automation, employing model distillation to meet low-latency requirements.
- Integrated Transformer architectures with prompt-engineering techniques to improve response-suggestion quality, reducing average ticket handling time by ~20%.
- Built churn-prediction models using gradient-boosted trees and introduced uplift modelling, boosting campaign ROI by ~25%.
- Implemented monitoring dashboards and alerting pipelines, enhancing issue-detection speed and supporting cross-team experimentation.

- Delivered ML projects across e-commerce, finance, and logistics, focusing on classical algorithms, forecasting, and automated reporting.
 - Established coding standards, experiment tracking, and comprehensive documentation, promoting reproducibility and collaborative development.
-

EDUCATION

M.Sc. Computer Science

Simulated University

B.Sc. Computer Science

Simulated University

PROJECTS

- Enterprise Q&A Assistant (RAG Platform): Multi-tenant retrieval-augmented generation system ingesting heterogeneous data sources, exposing a secure API for internal knowledge-base queries.
- LLM-Powered Customer-Support Agent: End-to-end pipeline from data preprocessing, fine-tuning (LoRA), prompt engineering, to production deployment on Kubernetes with auto-scaling.
- Churn Prediction & Uplift Modelling Suite: Integrated gradient-boosted trees with causal inference techniques, delivered as a reusable Python package.